

Time Series Analysis in R

Comprehensive Cheat Sheet

1 Data Handling and Exploration

Basic TS Functions

| | |
|----------------------------|--------------------------------|
| <code>ts()</code> | Create time series object |
| <code>c(ts_object)</code> | Convert TS to vector |
| <code>frequency(ts)</code> | Get observations per unit time |
| <code>cycle(ts)</code> | Get position in seasonal cycle |
| <code>time(ts)</code> | Extract time points |
| <code>window(ts)</code> | Extract subset of series |
| <code>log(ts)</code> | Log transformation |
| <code>as.ts()</code> | Convert object to time series |

Example

```
# Create time series of monthly data
my_ts <- ts(data, start=c(2020,1), frequency=12)

# Extract subset
subset_ts <- window(my_ts, start=c(2020,7), end=c(2021,6))
```

2 Visualization

Plotting Functions

| | |
|--------------------------------|--|
| <code>plot(ts)</code> | Basic time series plot |
| <code>lines(lowess())</code> | Add lowess smoother |
| <code>matplot()</code> | Plot matrix columns |
| <code>boxplot(ts~cycle)</code> | Boxplot by season |
| <code>layout(matrix())</code> | Arrange multiple plots |
| <code>ts.plot()</code> | Plot multiple series |
| <code>par(mfrow=c(r,c))</code> | Create plot grid |
| <code>pairs()</code> | Scatter plot matrix for multiple variables |

Example

```
# Multiple plots in one figure
par(mfrow=c(2,2))
plot(ts_data)
acf(ts_data)
pacf(ts_data)
cpgram(ts_data)

# Add smoother to plot
plot(ts_data)
lines(lowess(time(ts_data), ts_data, f=0.1), col="red", lwd=2)

# Seasonal plots
# Convert time series to matrix with one column per year
temp_matrix <- matrix(ts_data, nrow=12, byrow=FALSE)
matplot(temp_matrix, type="l")

# Boxplot by month/season
boxplot(ts_data ~ cycle(ts_data))
```

3 Decomposition and Smoothing

Decomposition Functions

| | |
|--|--|
| <code>decompose(ts)</code> | Classical decomposition using moving averages |
| <code>stl(ts, s.window)</code> | STL decomposition (Seasonal-Trend decomposition using Loess), more robust with irregular seasonality |
| <code>filter(ts, sides, filter)</code> | Moving average smoothing to remove noise and highlight trends |
| <code>lowess(x, y, f)</code> | LOWESS smoother for nonparametric trend fitting |

Example

```
# Classic decomposition
dcmp <- decompose(ts_data)
plot(dcmp)

# STL decomposition
stl_dcmp <- stl(ts_data, s.window="periodic")
plot(stl_dcmp)

# Moving average (centered, 3-point)
ma3 <- filter(ts_data, sides=2, rep(1,3)/3)

# Weighted moving average
wma <- filter(ts_data, sides=2, c(1,2,1)/4)
```

Moving Average Smoothers (Lab 3)

```
# Simple centered moving average
ma3 = filter(ts_data, sides=2, rep(1,3)/3)

# Weighted moving average
ma2.2 = filter(ts_data, sides=2, c(1,2,1)/4)

# Compare different smoothers
plot(ts_data, type="p")
lines(ma3, col="red")
lines(ma2.2, col="blue")
```

COSINOR Regression (Lab 3)

```
# Time is centered around mean time
wk = time(ts_data) - mean(time(ts_data))
wk2 = wk^2 # Quadratic term
wk3 = wk^3 # Cubic term

# Create Fourier terms
cs = cos(2*pi*wk)
sn = sin(2*pi*wk)

# Polynomial trend model
reg1 = lm(ts_data ~ wk + wk2 + wk3)

# Polynomial trend + seasonal components
reg2 = lm(ts_data ~ wk + wk2 + wk3 + cs + sn)

# Compare models
plot(ts_data)
lines(fitted(reg1), col="blue")
lines(fitted(reg2), col="red")
```

4 Correlation Analysis

Correlation Functions

| | |
|--|---|
| <code>lag(ts, k)</code> | Lag time series by k periods to examine relationships between current and past values |
| <code>acf(ts, lag.max)</code> | Autocorrelation function to identify overall correlation structure and seasonality |
| <code>pacf(ts, lag.max)</code> | Partial autocorrelation function to identify direct correlations (helps determine AR order) |
| <code>ccf(ts1, ts2)</code> | Cross-correlation function to analyze relationship between two time series |
| <code>ARMAacf(ar, ma, lag.max)</code> | Theoretical ACF for ARMA models to compare with empirical patterns |
| <code>ARMAacf(ar, ma, lag.max, pacf=TRUE)</code> | Theoretical PACF for ARMA models |

Example

```
# Plot autocorrelation function
acf(ts_data, lag.max=24)

# Compare empirical vs theoretical ACF
ar_model <- ar(ts_data)
theoretical_acf <- ARMAacf(ar=ar_model$ar, lag.max=20)

# Scatter plots of lagged values
plot(x=lag(ts_data, k=1), y=ts_data)
title(main=paste("r1", round(cor(ts_data[-1], ts_data[-length(ts_data)]), digits=3)))
```

5 Spectral Analysis

Spectral Functions

| | |
|-----------------------------|--|
| <code>spectrum(ts)</code> | Spectral density estimation to identify cyclical components |
| <code>cpgram(ts)</code> | Cumulative periodogram to check for white noise and hidden periodicities |
| <code>spec.pgram(ts)</code> | More control over periodogram calculation with options for smoothing |

Example

```
# Identify frequency components
spec <- spectrum(ts_data)

# Find dominant frequency
omega <- spec$freq[which.max(spec$spec)]
period <- 1/omega

# From Lab 7: Working with frequency components
# Calculate periodogram
spec_pgram <- spec.pgram(ts_data, spans=c(3,5),
                        taper=0.1, log="no")

# Frequency to period conversion
freqs <- spec_pgram$freq
periods <- 1/freqs[freqs > 0]
```

6 Stationarity

Stationarity Functions

| | |
|--------------------------------------|---|
| <code>diff(ts)</code> | First difference to remove trends and achieve stationarity |
| <code>diff(ts, lag=s)</code> | Seasonal difference to remove seasonal components |
| <code>Box.test(ts, lag, type)</code> | Box-Pierce or Ljung-Box test to test for autocorrelation in residuals |
| <code>adf.test()</code> | Augmented Dickey-Fuller test for stationarity (requires <code>tseries</code> package) |
| <code>kpss.test()</code> | KPSS test for trend stationarity (requires <code>tseries</code> package) |

Example

```
# Remove trend with differencing
ts_diff <- diff(ts_data)
plot(ts_diff)

# Remove seasonality (monthly data)
ts_sdifff <- diff(ts_data, lag=12)

# Seasonal differencing followed by regular differencing
ts_both_diff <- diff(diff(ts_data, lag=12))

# Test for stationarity
library(tseries)
adf_result <- adf.test(ts_data)
kpss_result <- kpss.test(ts_data)
```

7 AR, MA, and ARIMA Modeling

Model Fitting Functions

| | |
|---|---|
| <code>ar(ts, method, order.max)</code> | Fits autoregressive model (often uses Yule-Walker method) |
| <code>arima.sim(model, n)</code> | Simulates ARIMA process for generating data with known properties |
| <code>arima0(ts, order=c(p,d,q))</code> | Fits ARIMA model (faster version) for exploratory model fitting |
| <code>arima(ts, order=c(p,d,q))</code> | Fits ARIMA model with ML estimation for final model selection and inference |
| <code>arima(ts, order, seasonal)</code> | Fits seasonal ARIMA for data with seasonal patterns |
| <code>tsdiag(model)</code> | Diagnostic plots for time series models to check model adequacy |
| <code>residuals(model)</code> | Extracts residuals from a model for diagnostic checking |
| <code>AIC(model), BIC(model)</code> | Information criteria for model comparison and selection |
| <code>auto.arima()</code> | Automatic ARIMA model selection (requires <code>forecast</code> package) |

Example

```
# Fit ARIMA(1,1,1) model
model <- arima(ts_data, order=c(1,1,1))
summary(model)

# Seasonal ARIMA
sarima <- arima(ts_data, order=c(1,0,1),
                 seasonal=list(order=c(0,1,1), period=12))
```

Simulating ARMA Processes (Lab 4)

```
# Simulate AR(1) process
ar1_sim <- arima.sim(model=list(ar=0.7), n=100)

# Simulate AR(2) process
ar2_sim <- arima.sim(model=list(ar=c(0.7, -0.4)), n=100)

# Simulate ARMA(2,2) process
arma_sim <- arima.sim(model=list(ar=c(0.7, -0.4),
                                   ma=c(-0.2, 0.25)),
                      n=100,
                      sd=sqrt(2))

# Simulate with different error distribution
arma_t_sim <- arima.sim(model=list(ar=c(0.7, -0.4),
                                   ma=c(-0.2, 0.25)),
                      n=100,
                      rand.gen=function(n,...) {
                        return(sqrt(0.2)*rt(n, df=5))
                      })
```

AR(2) Stationarity Conditions (Lab 4)

For an AR(2) process to be stationary, the parameters ϕ_1 and ϕ_2 must satisfy:

- $\phi_2 > -1$
- $\phi_2 + \phi_1 < 1$
- $\phi_2 - \phi_1 < 1$

```
# Check stationarity by analyzing roots
# For AR(2): (1-B-B)Y = e
# Roots must be outside unit circle
roots <- polyroot(c(1, -0.7, 0.4))
abs(roots) # Should be > 1 for stationarity
```

8 Forecasting

Forecasting Functions

| | |
|--------------------------------------|---|
| <code>predict(model, n.ahead)</code> | Basic forecasting from ARIMA models |
| <code>predict(model, newdata)</code> | Forecasting from regression models with predictors |
| <code>forecast()</code> | Enhanced forecasting with visualizations (requires <code>forecast</code> package) |

Example

```
# Forecast next 12 periods
fore <- predict(arima_model, n.ahead=12)

# Plot with prediction intervals
ts.plot(ts_data, fore$pred,
        fore$pred + 2*fore$se,
        fore$pred - 2*fore$se,
        gpars=list(col=c(1,2,4,4)))

# Using forecast package
library(forecast)
model_fc <- forecast(arima_model, h=12)
plot(model_fc)
```

9 Regression Models for Time Series

Regression Functions

| | |
|--------------------------------------|--|
| <code>lm(y ~trend + seasonal)</code> | Linear regression with trend and seasonality to model deterministic patterns |
| <code>lm(y ~cbind(x1, x2))</code> | Regression with multiple predictors for models with covariates |
| <code>cos(2*pi*time()), sin()</code> | Fourier terms for modeling seasonality in regression |
| <code>summary(aov(model))</code> | ANOVA table for regression model |

Example

```
# Trend with harmonic seasonality
t <- time(ts_data)
s <- sin(2*pi*t)
c <- cos(2*pi*t)
reg <- lm(ts_data ~ t + s + c)

# Polynomial trend terms
t2 <- t^2
t3 <- t^3
reg_poly <- lm(ts_data ~ t + t2 + t3 + s + c)

# Multiple harmonics (from Lab 3)
s1 <- sin(2*pi*t)
c1 <- cos(2*pi*t)
s2 <- sin(4*pi*t) # Second harmonic (twice the frequency)
c2 <- cos(4*pi*t)
reg_harm <- lm(ts_data ~ t + s1 + c1 + s2 + c2)

# ANOVA analysis
summary(aov(reg_harm))
```

10 Multivariate Time Series

Multivariate Functions

| | |
|-----------------------------------|---|
| <code>VAR(ts_data, p)</code> | Vector Autoregression for modeling relationships between multiple time series |
| <code>VARselect(ts_data)</code> | Helps select VAR order to determine optimal lag order in VAR models |
| <code>irf(var_model)</code> | Impulse response function analysis |
| <code>fevd(var_model)</code> | Forecast error variance decomposition |
| <code>causality(var_model)</code> | Granger causality tests |

VAR Modeling Example (Lab 7)

```
library(vars)

# Create multivariate time series
mts_data <- ts(cbind(series1, series2), frequency=12)

# Select optimal lag order
lag_selection <- VARselect(mts_data, lag.max=10,
                           type="const")
optimal_p <- lag_selection$selection[["SC"]] # Using Schwarz criterion

# Fit VAR model
var_model <- VAR(mts_data, p=optimal_p, type="const")
summary(var_model)

# Analyze impulse response
irf_result <- irf(var_model, n.ahead=12)
plot(irf_result)

# Forecast error variance decomposition
fevd_result <- fevd(var_model, n.ahead=10)
plot(fevd_result)

# Granger causality
causality(var_model, cause="series1")
```

11 ACF/PACF Pattern Recognition

Model Identification Patterns

| Process | ACF | PACF |
|-------------|--------------------------|--------------------------|
| AR(p) | Tails off gradually | Cuts off after lag p |
| MA(q) | Cuts off after lag q | Tails off gradually |
| ARMA(p,q) | Tails off gradually | Tails off gradually |
| White Noise | No significant spikes | No significant spikes |
| Seasonal | Spikes at lags s, 2s, 3s | Spikes at lags s, 2s, 3s |

12 Model Diagnostic Steps

1. **Plot residuals** - should resemble white noise
2. **ACF of residuals** - no significant autocorrelation
3. **Ljung-Box test** - p-values should exceed significance level
4. **QQ-plot** - check normality assumption

5. Cumulative periodogram - should follow diagonal line

Diagnostic Checks (Lab 5)

```
# Fit model
model <- arima(ts_data, order=c(1,0,1))

# Run diagnostics
tsdiag(model)

# Manual diagnostics
par(mfrow=c(2,2))
# Plot residuals
plot(residuals(model), main="Residuals")
# ACF of residuals
acf(residuals(model), main="ACF of Residuals")
# Ljung-Box test p-values
lb_pvalues <- sapply(1:20, function(i)
  Box.test(residuals(model), lag=i, type="Ljung-Box")$p.value)
plot(lb_pvalues, main="Ljung-Box p-values",
  ylab="p-value", xlab="lag")
abline(h=0.05, col="red", lty=2)
# QQ plot
qqnorm(residuals(model))
qqline(residuals(model), col="red")

# Cumulative periodogram
cpgram(residuals(model))
```

13 ARIMA Modeling Workflow

1. **Plot data** - identify patterns and anomalies
2. **Transform if needed** - typically log transformation
3. **Difference until stationary** - regular/seasonal
4. **Examine ACF/PACF** - identify potential orders
5. **Fit candidate models** - try several specifications
6. **Compare using AIC/BIC** - select best model
7. **Check diagnostics** - validate residual behavior
8. **Forecast** - generate predictions with intervals

```

# Systematic model comparison
models <- list()
aic_values <- matrix(NA, nrow=3, ncol=3)
bic_values <- matrix(NA, nrow=3, ncol=3)

for (p in 0:2) {
  for (q in 0:2) {
    model_name <- paste("ARIMA(", p, ",0,", q, ")", sep="")
    models[[model_name]] <- arima(ts_data, order=c(p,0,q))
    aic_values[p+1, q+1] <- AIC(models[[model_name]])
    bic_values[p+1, q+1] <- BIC(models[[model_name]])
  }
}

# Find best model by AIC
min_aic <- which(aic_values == min(aic_values), arr.ind=TRUE)
best_p_aic <- min_aic[1] - 1
best_q_aic <- min_aic[2] - 1
cat("Best model by AIC: ARIMA(", best_p_aic, ",0,", best_q_aic, ")\n", sep="")

# Find best model by BIC
min_bic <- which(bic_values == min(bic_values), arr.ind=TRUE)
best_p_bic <- min_bic[1] - 1
best_q_bic <- min_bic[2] - 1
cat("Best model by BIC: ARIMA(", best_p_bic, ",0,", best_q_bic, ")\n", sep="")

```

14 Key Exam Insights

- **Stationarity** is fundamental - constant mean, variance, and autocorrelation
- **Transformation sequence:** log → seasonal differencing → regular differencing
- **Model parsimony:** Simpler models often forecast better
- **Residual analysis:** The true test of any model
- **Seasonal patterns:** Use seasonal ARIMA or Fourier terms
- **Combined approaches:** Regression for deterministic components, ARIMA for residuals
- **Spectral analysis:** Identify hidden periodicities
- **Smoothing techniques:** Extract trends flexibly
- **Multivariate analysis:** Use VAR for interdependent series

15 Practical Tips for Bike Share Analysis

Data Preparation and Model Selection

```
# Step 1: Load data
load("sharedbikes.RData")
load("sharedbikes_forecast.RData")

# Step 2: Linear model for trend/seasonality
mod <- lm(casual ~ t + s1 + c1 + s2 + c2 + s7_1 + c7_1 + s7_2 + c7_2, data=dat)

# Step 3: Examine residuals
dat$res <- ts(residuals(mod), start=2018, frequency=365)
plot(dat$res)
acf(dat$res)
pacf(dat$res)

# Step 4: Test for stationarity with differencing
diff_res <- diff(dat$res)
plot(diff_res)
acf(diff_res)
cpgram(diff_res)

# Step 5: Find best ARIMA model for residuals
# Initialize matrix for AIC values
p_values <- 0:6
q_values <- 0:6
aic_matrix <- matrix(NA, nrow=length(p_values), ncol=length(q_values))
rownames(aic_matrix) <- paste("p_", p_values)
colnames(aic_matrix) <- paste("q_", q_values)

# Fit ARIMA(p,1,q) models and store AIC values
for (i in 1:length(p_values)) {
  for (j in 1:length(q_values)) {
    p <- p_values[i]
    q <- q_values[j]

    tryCatch({
      model <- arima0(dat$res, order=c(p, 1, q))
      aic_matrix[i, j] <- model$aic
    }, error = function(e) {
      aic_matrix[i, j] <- NA
    })
  }
}

# Find the model with the lowest AIC
min_aic <- min(aic_matrix, na.rm=TRUE)
min_indices <- which(aic_matrix == min_aic, arr.ind=TRUE)
best_p <- p_values[min_indices[1]]
best_q <- q_values[min_indices[2]]

# Step 6: Fit best model and check diagnostics
best_model <- arima(dat$res, order=c(best_p, 1, best_q))
tsdiag(best_model)
```

```

# Step 7: Generate forecasts
# Predict from linear model
linear_forecast <- predict(mod, newdata=newdat)

# Forecast residuals with ARIMA
arima_forecast <- predict(best_model, n.ahead=60)

# Combine forecasts
total_forecast <- linear_forecast + arima_forecast$pred

# Create confidence intervals
upper_ci <- total_forecast + 1.96 * arima_forecast$se
lower_ci <- total_forecast - 1.96 * arima_forecast$se

# Step 8: Plot forecasts
# Create date sequence for forecast period
forecast_dates <- seq.Date(from=max(dat$dteday) + 1,
                           by="day", length.out=60)

# Plot original data and forecasts
plot(dat$dteday, dat$casual,
     type="l",
     xlim=c(min(dat$dteday), max(forecast_dates)),
     ylim=c(min(c(dat$casual, lower_ci), max(c(dat$casual, upper_ci))),
     xlab="Date", ylab="Number_of_Casual_Bike_Hires",
     main="Casual_Bike_Rentals_with_60-day_Forecast")

# Add forecast and prediction intervals
lines(forecast_dates, total_forecast, col="blue", lwd=2)
lines(forecast_dates, upper_ci, col="red", lty=2)
lines(forecast_dates, lower_ci, col="red", lty=2)

# Add legend
legend("topleft",
     legend=c("Observed", "Forecast", "95%_Prediction_Interval"),
     col=c("black", "blue", "red"),
     lty=c(1, 1, 2),
     lwd=c(1, 2, 1))

```